

Case 5.4 Build Regression using Regression Node (Using R)

5.4.1 Purpose of the Study

To understand how to build a regression model using R.

5.4.2 Data

The data set used in this study is the same data used in Association for Computing Machinery's (ACM) 1998 KDD-Cup competition (<http://kdd.ics.uci.edu>). The original data is much more complicated and the version used is here from SAS Institute modified version. The response rate for this donation campaign is about 5% although you can see the response rate in this subsample is 50%. This means the prior probability is 0.05.

```
require(readxl)
cup98LRN <- as.data.frame(read_excel("F:\\R\\pmad_pva.xlsx"))
cup98LRN <- cup98LRN[!is.na(cup98LRN$TargetD),]
```

A great aspect of the caret package is that it handles a good portion of data preparation for us. Here we use the (preProcess) function to

- remove zero or near zero variance columns
- remove columns that are highly correlated with other columns
- Box-Cox transform the columns where appropriate
- center and scale the columns to have mean zero variance one
- impute missing values with median values
- apply the spatial sign transform

```
y <- cup98LRN$TargetD
X <- cup98LRN[,-c(1:3)]
require(caret)
```

```
preProcess.X <- preProcess(X, method = c(
  "zv", "nzv", "corr", "BoxCox", "center", "scale", "medianImpute", "spatialSign"
),
  outcome = y
)
X <- predict(preProcess.X, X)
```

There are several categorical variables. Here we bin each categorical vector into two groups; one group is a bin of the levels with smaller conditional averages of TargetD, the other group is a bin of the levels with larger conditional averages of TargetD. The bins were constructed so that the smaller/larger partition gives them a split as close to 50-50 as possible. The bins are represented as binary numeric

Case 5.4 Build Regression using Regression Node (Using R)

vectors. Notice that after binning all columns are numeric vectors; many model functions complete their computations faster with numeric variables.

```
X$StatusCat96NK <- as.numeric(X$StatusCat96NK == "S") BIN.DemCluster <- c(
  "04","08","09","16","19","20","23","25","26","27",
  "28","30","32","33","36","38","39","40","41","43",
  "45","46","47","48","49","51","52","53"
)
X$DemCluster <- as.numeric(X$DemCluster %in% BIN.DemCluster) X$DemGender <-
as.numeric(X$DemGender == "F")
X$DemHomeOwner <- as.numeric(X$DemHomeOwner == "H")
```

5.4.3 Regression Modeling in R

```
list.model <- list() list.model[["lm"]] <- train(
  y = y, x = X,
  trControl = trainControl(method = "repeatedcv", number = 5, repeats = 5), method = "lm"
)
for(j in c("forward", "both", "backward")){ list.model[[j]] <- train(
  y = y, x = X,
  trControl = trainControl(method = "repeatedcv", number = 5, repeats = 5), method =
  "lmStepAIC",
  trace = 0, direction = j
)
}
```

Case 5.4 Build Regression using Regression Node (Using R)

Lets look at the coefficients.

```
sapply(X = list.model,FUN = function(x) coef(x$finalModel))
```

```
## $lm
```

```
##      (Intercept)      GiftCnt36      GiftCntCard36      GiftCntCardAll
```

Case 5.4 Build Regression using Regression Node (Using R)

```
##      17.7457247      1.7967810      -1.4469884      -1.3292366
##      GiftAvgLast      GiftAvg36      GiftAvgAll      GiftTimeLast
##      31.1333606      23.8100025      -3.8781869      0.6333394
##      GiftTimeFirst      PromCnt12      PromCnt36      PromCntAll
##      -4.2509375      3.0695342      -4.7773528      3.2241325
##      PromCntCard12      PromCntCard36      StatusCat96NK      StatusCatStarAll
##      -0.8038078      1.4134616      -0.5439398      0.5632963
##      DemCluster      DemAge      DemGender      DemHomeOwner
##      -0.6793548      0.3396813      -0.6120642      -0.6880468
##      DemMedHomeValueDemPctVeterans      DemMedIncome
##      -0.2974096      0.7663830      0.1339993
##
## $forward
##      (Intercept)      GiftCnt36      GiftCntCard36      GiftCntCardAll
##      17.7457247      1.7967810      -1.4469884      -1.3292366
##      GiftAvgLast      GiftAvg36      GiftAvgAll      GiftTimeLast
##      31.1333606      23.8100025      -3.8781869      0.6333394
##      GiftTimeFirst      PromCnt12      PromCnt36      PromCntAll
##      -4.2509375      3.0695342      -4.7773528      3.2241325
##      PromCntCard12      PromCntCard36      StatusCat96NK      StatusCatStarAll
##      -0.8038078      1.4134616      -0.5439398      0.5632963
##      DemCluster      DemAge      DemGender      DemHomeOwner
##      -0.6793548      0.3396813      -0.6120642      -0.6880468
##      DemMedHomeValueDemPctVeterans      DemMedIncome
##      -0.2974096      0.7663830      0.1339993
##
## $both
##      (Intercept)      GiftCnt36      GiftCntCard36      GiftAvgLast      GiftAvg36
##      17.7660595      1.6435803      -1.6973733      31.1680536      23.6571237
##      GiftAvgAll      GiftTimeFirst      PromCnt12      PromCnt36      PromCntAll
##      -3.7508644      -4.5883317      2.0617651      -4.2861356      2.9553127
##      PromCntCard36      StatusCat96NK      DemCluster      DemGender      DemHomeOwner
##      1.2345969      -0.6015682      -0.6397340      -0.6221611      -0.6837266
##      DemPctVeterans
##      0.8079780
##
## $backward
##      (Intercept)      GiftCnt36      GiftCntCard36      GiftAvgLast      GiftAvg36
##      17.7660595      1.6435803      -1.6973733      31.1680536      23.6571237
##      GiftAvgAll      GiftTimeFirst      PromCnt12      PromCnt36      PromCntAll
##      -3.7508644      -4.5883317      2.0617651      -4.2861356      2.9553127
##      PromCntCard36      StatusCat96NK      DemCluster      DemGender      DemHomeOwner
##      1.2345969      -0.6015682      -0.6397340      -0.6221611      -0.6837266
##      DemPctVeterans
##      0.8079780
```

Now let's look at cross-validation statistics.

Case 5.4 Build Regression using Regression Node (Using R)

```
sapply(X = list.model,FUN = getTrainPerf)

##      lm      forward both      backward
## TrainRMSE 9.026997      9.01431      9.051548      9.065909
## TrainRsquared 0.4738884 0.4755433      0.4718298      0.4699284
## TrainMAE 4.691627      4.693219      4.697055      4.698826
## method      "lm"      "lmStepAIC" "lmStepAIC" "lmStepAIC"
```

There appears to be two models, the full model and a reduced model. Let's run an ANOVA test to see if the full model is sufficiently better than the reduced model.

```
reduced.model <- list.model[["backward"]]$finalModel
full.model <- list.model[["lm"]]$finalModel anova(reduced.model,full.model)

## Analysis of Variance Table
##
## Model 1: .outcome ~ GiftCnt36 + GiftCntCard36 + GiftAvgLast + GiftAvg36 +
##      GiftAvgAll + GiftTimeFirst + PromCnt12 + PromCnt36 + PromCntAll +
##      PromCntCard36 + StatusCat96NK + DemCluster + DemGender +
##      DemHomeOwner + DemPctVeterans
## Model 2: .outcome ~ GiftCnt36 + GiftCntCard36 + GiftCntCardAll + GiftAvgLast +
##      GiftAvg36 + GiftAvgAll + GiftTimeLast + GiftTimeFirst + PromCnt12 +
##      PromCnt36 + PromCntAll + PromCntCard12 + PromCntCard36 +
##      StatusCat96NK + StatusCatStarAll + DemCluster + DemAge + DemGender +
##      DemHomeOwner + DemMedHomeValue + DemPctVeterans +
##      DemMedIncome
##      Res.Df      RSS Df Sum of Sq      F Pr(>F)
## 1      4827 398497
## 2      4820 398192      305.32 0.528 0.814
##      7
```

It looks like the reduced model is just as good as the full model. Since they are equally good, the simpler one is better; let's check it out.

Case 5.4 Build Regression using Regression Node (Using R)

[illegible]

Case 5.4 Build Regression using Regression Node (Using R)

```
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   17.7661    0.3027  58.690 < 2e-16 ***
## GiftCnt36      1.6436    1.0340   1.589 0.112016
## GiftCntCard36 -1.6974    0.9727  -1.745 0.081053 .
## GiftAvgLast   31.1681    1.6170  19.275 < 2e-16 ***
## GiftAvg36     23.6571    1.6041  14.748 < 2e-16 ***
## GiftAvgAll    -3.7509    1.0819  -3.467 0.000531 ***
## GiftTimeFirst -4.5883    1.6061  -2.857 0.004296 **
## PromCnt12      2.0618    1.0793   1.910 0.056165 .
## PromCnt36     -4.2861    1.5779  -2.716 0.006624 **
## PromCntAll     2.9553    2.0924   1.412 0.157896
## PromCntCard36  1.2346    0.8348   1.479 0.139246
## StatusCat96NK -0.6016    0.3681  -1.634 0.102276
## DemCluster    -0.6397    0.2693  -2.376 0.017545 *
## DemGender     -0.6222    0.2625  -2.370 0.017839 *
## DemHomeOwne  -0.6837    0.2651  -2.580 0.009922 **
## r
## DemPctVeterans 0.8080    0.5125   1.577 0.114944
## ---
```

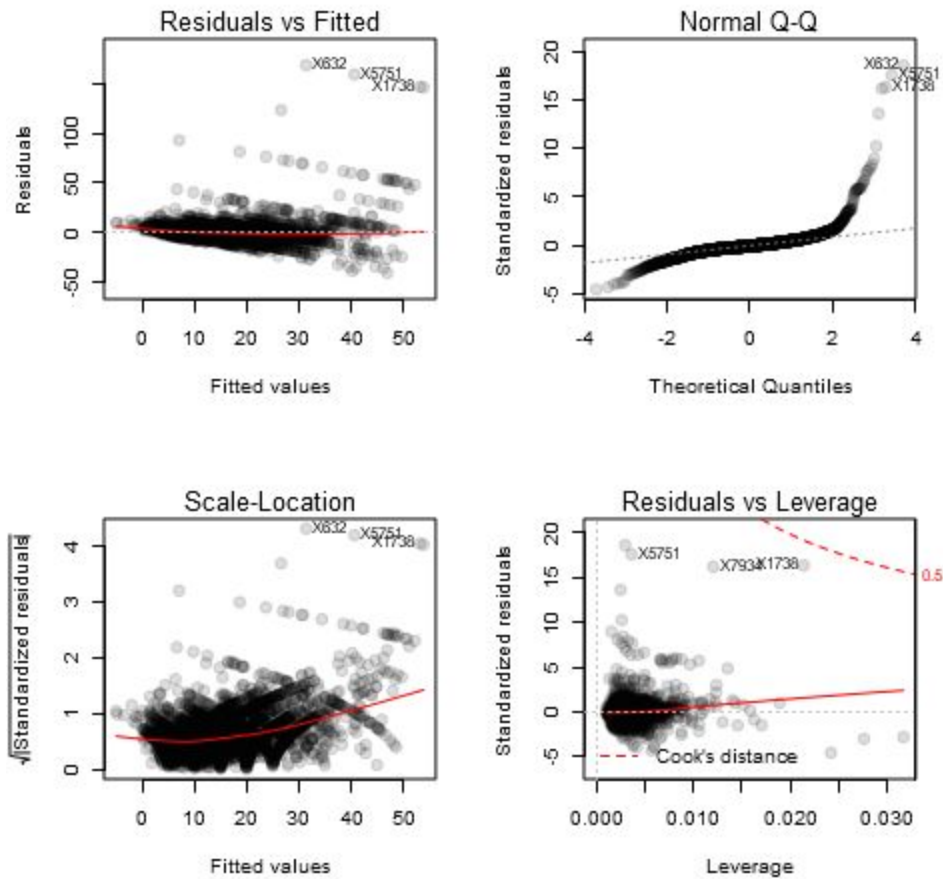
Case 5.4 Build Regression using Regression Node (Using R)

```
## GiftAvg36      1 18867 18867 228.5421 < 2.2e-16 ***
## GiftAvgAll     1   30    30   0.3663  0.54503
## GiftTimeFirst  1 3028  3028  36.6813 1.497e-09 ***
## PromCnt12      1   28    28   0.3439  0.55759
## PromCnt36      1  293   293   3.5517  0.05954 .
1
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
plot(reduced.model, pch = 19, col = "#00000022")
par(mfrow = c(2,2))
```

```
## PromCntAll      97    97   1.1734  0.27876
## PromCntCard36   1  144   144   1.7502  0.18591
## StatusCat96NK   1  218   218   2.6377  0.10442
## DemCluster      1  336   336   4.0733  0.04362 *
## DemGender       1  460   460   5.5690  0.01832 *
## DemHomeOwne     1  497   497   6.0163  0.01421 *
r
## DemPctVeterans  1  205   205   2.4858  0.11494
## Residuals      4827 398497   83
## ---
```


Case 5.4 Build Regression using Regression Node (Using R)



```
dev.off()

## windows
##      2
```

Appendix 1 Discussion Questions

What are the optimal parameter settings for entry into a forward selection model?

Case 5.4 Build Regression using Regression Node (Using R)

1. What are the optimal parameter settings for exit from a backward elimination model?
2. Why we prefer to use validation error on selecting the best model?